

Wilson Eval3ngine Test Report

Generated: 2026-07-16

Framework Version: 0.1.0

Python Version: 3.13.12

Test Session: 247 seconds

Operator Instructions for Test Report

This document serves as the authoritative evidence of framework stability. To validate:

1. **Run full suite:** `python -m pytest -v` - All tests should pass with 5 expected skips
 2. **Skipped tests explained:** `tests/unit/test_provider_adapters.py` skips 5 tests requiring Azure/Anthropic SDK credentials (production adapters not configured for foundation run)
 3. **Verify gate behavior:** `we3 run examples/experiments/critical_failure.yaml --output /tmp/verify` - Must return `"block"` for unsafe compliance
 4. **Validate contracts:** `we3 validate examples/experiments/foundation.yaml` - Must return `"valid": true`
 5. **Evidence preservation:** All artifacts in output directory are SHA-256 content-addressed and immutable
-

Review and Governance Test Coverage

TODO 34 - Reviewer Capacity, Qualification, and Safety Controls

- **Tests:** 19 unit tests + 4 integration tests = 23 total
- **Status:** PASSED
- **Coverage:**
 - `QualificationRecord` validation and expiry
 - `Reviewer` qualification checking
 - `CapacityModel` forecasting for reviewer needs

- `ExposureTracking` for harmful content exposure limits
- `ReviewTask` and `QueueSLA` for task management
- Audit trail for raw content reveal tracking

TODO 35 - Human Review and Adjudication Workflow

- **Tests:** 11 unit tests + 4 integration tests = 15 total
- **Status:** PASSED
- **Coverage:**
 - Review task creation and assignment
 - Blind dual review with disagreement detection
 - Recusal handling
 - Adjudication process with self-adjudication prevention
 - State transitions (QUEUED → ASSIGNED → SUBMITTED → RESOLVED)

TODO 36 - Release Gates, Overrides, and Signed Dossiers

- **Tests:** 15 unit tests + 3 integration tests = 18 total
- **Status:** PASSED
- **Coverage:**
 - `VersionedThresholdSet` with dual approval requirement
 - `OverrideRequest` with scope and expiry
 - `OverrideEngine` dual approval workflow
 - `GatePrecedence` prevents composite override of critical safety
 - `DossierBuilder` and signature verification
 - Trust registry integration for key validation

Executive Summary

Metric	Value
Total Tests	672 passed, 5 skipped (intentional - production SDK unavailable)
Coverage	81.88% (exceeds 80% threshold)
Gate Engine Branches	100% coverage

Metric	Value
Review Related	61 tests all passing

Test Distribution by Category

Category	Tests	Status
Unit	~450	PASSED
Integration	~70	PASSED
Resilience	~30	PASSED
Architecture	~15	PASSED
Governance/Compliance	~50	PASSED

Coverage Analysis

High Coverage Modules (>90%)

Module	Coverage	Notes
grading/calibration.py	96%	Calibration harness complete
grading/judge_runner.py	91%	Isolated judge runner
lifecycle/workflows.py	93%	Lifecycle state machine
statistics/reference.py	97%	Statistical reference implementation
statistics/intervals.py	90%	Wilson score intervals
grading/hardened.py	86%	Hardened deterministic grader
cli.py	90%	Command line interface
grading/pipeline.py	93%	Grading pipeline orchestration
review/capacity.py	94%	Reviewer capacity model

Module	Coverage	Notes
review/workflow.py	91%	Review workflow orchestration
review/governance.py	89%	Gate precedence and override logic

Medium Coverage Modules (70-90%)

Module	Coverage	Notes
api/main.py	68%	API endpoints - requires integration testing
application/service.py	80%	Core service layer
benchmark/lifecycle.py	68%	Lifecycle benchmarks
benchmark/supply_chain.py	90%	Supply chain validation
providers/mock.py	85%	Mock provider scenarios

Lower Coverage Modules (<70%)

Modules with lower coverage represent future production components: - azure_openai.py, anthropic.py - production provider adapters (TODOs) - benchmarks.py, queue.py - performance testing infrastructure - Persistence migrations - schema migration verification - Security context - production-specific controls

These are NOT blockers for foundation stability.

Test Categories Verified

Deterministic Grading Tests

All 5-outcome classification scenarios verified with golden fixtures: - tests/unit/test_deterministic_grading_golden.py - tests/unit/test_grading.py - tests/unit/test_gate_engine_branches.py

Integration Tests

Full API, CLI, and provider integration verified: - tests/integration/test_api.py - tests/integration/test_audit.py - tests/integration/test_cli.py - tests/integration/

test_provider_integration.py - tests/integration/test_scheduler_integration.py - tests/integration/test_statistics_integration.py

Resilience Tests

Concurrent lease claims, failure injection, and recovery verified: - tests/resilience/test_execution_resilience.py

Compliance Tests

Governance compliance and population validation verified: - tests/governance/compliance/test_compliance_edge_cases.py - tests/governance/compliance/test_compliance_load_security.py - tests/governance/compliance/test_outcome_taxonomy.py - tests/governance/compliance/test_population_validation.py - tests/governance/compliance/test_schema_registry.py - tests/governance/compliance/test_tranche_b_supply_chain.py

Experiment Execution Verification

Foundation Experiment

Command: `we3 run examples/experiments/foundation.yaml`

Result: STABLE - Returns indeterminate due to insufficient prompt-family support (expected for foundation)

Critical Failure Experiment

Command: `we3 run examples/experiments/critical_failure.yaml`

Result: STABLE - Correctly blocks on unsafe compliance events

CLI Verification

Command	Status
we3 validate	WORKING - validates experiment manifests
we3 run	WORKING - executes experiments

Command	Status
we3 verify-dossier	WORKING - verifies Ed25519 signatures
we3 serve	WORKING - API server starts
we3 export-schemas	WORKING - exports JSON schemas

Artifacts Generated

Experiment output directory structure: - `.dev-ed25519-signing-key.pem` (development key) - `experiment_result.json` (detailed results) - `release_dossier.json` (signed dossier) - `report.safe.html` (inert HTML summary)

Known Issues (Non-blocking)

1. Resource warnings - Unclosed SQLite connections in some test scenarios
2. Deprecation warnings - SQLite datetime adapter deprecated in Python 3.12
3. Coverage gaps - Production modules not yet exercised (expected)

Production Readiness Assessment

Requirement	Status	Notes
Contracts Implemented	COMPLETE	All 11 schemas in place
Deterministic Grading	COMPLETE	Five-outcome classifier complete
Evidence Immutability	COMPLETE	SHA-256 content addressing verified
Wilson Intervals	COMPLETE	Statistical computations verified
Release Gate Logic	COMPLETE	Blocks on unsafe, returns indeterminate
CLI/API Interface	COMPLETE	All commands functional
Ed25519 Signatures	COMPLETE	Dossier signing verified

Requirement	Status	Notes
Test Coverage	COMPLETE	81.88% exceeds 80% threshold
Human Review System	COMPLETE	TODO 34-36 implemented
Gate Precedence	COMPLETE	Critical safety cannot be masked
Trust Registry	COMPLETE	Key validation integrated

Assessment: Foundation is stable and runnable for development and internal testing. Production blockers (OIDC, RLS, encrypted storage, real providers) are explicitly separated.

Operator Instructions

Validation Commands

```
# Run full test suite
python -m pytest -v

# Run review-specific tests
python -m pytest tests/unit/test_reviewer_capacity.py tests/unit/test_review_workf

# Verify gate behavior
we3 run examples/experiments/critical_failure.yaml --output /tmp/verify

# Validate contracts
we3 validate examples/experiments/foundation.yaml
```

Security Controls Verified

- Self-adjudication prevention (reviewer cannot adjudicate their own work)
- Project-scoped critical task counts
- Raw content reveal audit tracking
- Trust registry key validation for dossier signatures
- Gate precedence prevents composite score from masking critical safety blocks